

The native data-centric recipe lifts every model — in detection and OBB

Mean $\pm$ std over 5 disjoint folds, 1,343-image merged pool · no external data, no new labels · detection best = champ v11n; OBB best = osall v11n
OBB folds use a separate annotation export — compare within task (same-fold control: task mode is a wash, det 0.802 vs OBB 0.808 mAP)
■ base (640 px) ■ champ (1280 px + 3× oversample + scale 0.9) ■ osall (all-defect 3× oversample)

