


768 px native retrain (circled) is the deployment point — FP16 export costs only -0.009 mAP@0.5

Resolution ladder, YOLOv11n champ recipe on the clean benchmark · retraining natively at 768 beats downscaled 1280 inference at 2.8× less compute



* fps measured on-device on a Jetson Orin Nano (PyTorch fp32, batch-1, end-to-end incl. pre/post; hardware_testing/jetson_fps_results.json) — real, not a projection. Not comparable to the original (Maxwell) Jetson Nano.